

Improving model performance: Model Selection, Parameter Tuning and Feature Engineering

B. Michel

Ecole Centrale de Nantes

Statistical Learning


Outline

- 1 **Overfitting and Bias Variance Tradeoff**
- 2 Estimating the generalization error by cross validation
- 3 Regularization, model selection and penalization
- 4 Feature engineering



Overfitting

- Statistics is not just interpolating between observation points !
- Overfitting occurs when the learning rule too closely corresponds to a particular set of data, and may therefore fail to generalize on future observations.

Regression:






predictor too inflexible:
cannot capture pattern



predictor too flexible:
fits noise in the data

Classification:



Copyright © 2014 Victor Lavyenko

Overfitting


- Statistics is not just interpolating between observation points !
- Overfitting occurs when the learning rule too closely corresponds to a particular set of data, and may therefore fail to generalize on future observations.
- An overfitted model is a statistical model that contains more parameters than can be learned (estimated) by the data.
- Examples :
 - ▶ In regression : e.g. a linear model with too many variables (features).
 - ▶ In classification : e.g. a k-NN with k too small.
 - ▶ In density estimation : an histogram with too many bins.
 - ▶ Kernel methods : bandwidth too small
 - ▶ ...
- High dimensional statistics : n is of the order (or smaller than) the number of features p .

Bias Variance Tradeoff

- loss ℓ : quadratic, logistic, hinge loss
- Risk or Generalization error of a learning rule $\hat{f} : \mathbb{E}\ell(Y, \hat{f}(X))$
- Typical decomposition of the generalization error :

Generalization error = (or \leq) Approximation error + Estimation error


- Estimation error : random term \approx variance term in statistics.
- Approximation error : deterministic term \approx bias term in statistics.



[Daniel Saunders]

Training error versus test error

- How can we **estimate** the generalization error from the data ?
- The **learning error** is the average error that results from estimating the performance of a learning rule using the sample **already used** to fit the learning rule.
- The **test error** is the average error that results from estimating the performance of a learning rule using a **different** sample.
- The **learning error** can dramatically **underestimate** the generalization error :



Question: Why does the learning error may fail to estimate the generalization error ?

Outline

- 1 Overfitting and Bias Variance Tradeoff
- 2 Estimating the generalization error by cross validation**
- 3 Regularization, model selection and penalization
- 4 Feature engineering

Splitting

- Sample $(X_i, Y_i)_{i=1, \dots, n}$, loss function ℓ .
- Split the data at random :
 - ▶ I_{train} : observations in the train set ;
 - ▶ I_{test} : observations in the test set.
- Learn a learning rule \hat{f} on I_{train} .
E.g Empirical Risk Minimization (ERM) on S with loss ℓ :

$$\hat{f} = \operatorname{argmin}_{f \in S} \frac{1}{|I_{\text{train}}|} \sum_{i \in I_{\text{train}}} \ell(Y_i, f(X_i))$$

- Compute the test error on I_{test} :

$$\text{Err}_{\text{test}} = \frac{1}{|I_{\text{test}}|} \sum_{i \in I_{\text{test}}} \ell(Y_i, \hat{f}(X_i))$$


- Split at random N times and average the N test errors.

Splitting

- For each splitting we compute a learning rule : high computational cost.
- The train sample is smaller than the complete sample : it overestimates the generalization error.
- Alternative : Cross validation methods.

Cross validation : k-fold cross validation

- k-fold validation: pick a random partition $I_1, \dots, I_v, \dots, I_k$ of $\{1, \dots, n\}$ where $|I_v| \sim n/k$.




[Gaiffas]

For each fold v :

- $I_{\text{test},v} = I_v$,
 - $I_{\text{train},v} = \{1, \dots, n\} \setminus I_v$,
 - \hat{f}_v : learning rule fitted (trained) on $I_{\text{train},v}$.
- Compute
- $$\text{Err} = \frac{1}{k} \sum_{v=1 \dots k} \frac{1}{|I_v|} \sum_{i \in I_v} \ell(Y_i, \hat{f}_v(X_i)).$$
- Only k learning rules are fitted with this approach.

Cross validation : Leave one out

- Leave one out (Loo) : $k = n$



[From : The Elements of Stat Learning]

- $I_{\text{test},v} = \{v\},$
- $I_{\text{train},v} = \{1, \dots, n\} \setminus \{v\},$

Cross validation

- Cross validation aims at estimating the generalization error :
 - ▶ Cross validation methods are estimation procedures, they have a bias and a variance.
 - ▶ In practice : take $k \approx 5 - 10$.
 - ▶ Cross validation with sklearn :
- Cross validation can be used for
 - ▶ tuning a parameter, e.g. for regularization.
 - ▶ for selecting models in a given collection.

In both cases: **after tuning the parameters and selecting a model, you should fit the chosen model on the complete dataset (not only on the train dataset).**

Outline

- 1 Overfitting and Bias Variance Tradeoff
- 2 Estimating the generalization error by cross validation
- 3 Regularization, model selection and penalization**
- 4 Feature engineering

Linear regression in high dimension

- Regression, design matrix X of size $n \times p$, target Y .
- Linear predictors of Y : $f_{\beta}(x) = x'\beta$.
- Quadratic loss : $\hat{\beta} = (X'X)^{-1}X'Y$.
- High dim: $n \sim p$ or $n \lesssim p$: $X'X$ is not invertible.
- Ridge regularization :

$$\hat{\beta}_{\lambda}^{\text{ridge}} = \underset{\beta \in \mathbb{R}^p}{\operatorname{argmin}} \|Y - X\beta\|^2 + \lambda \sum_{j=1}^p \beta_j^2.$$

and the solution is

$$\hat{\beta}_{\lambda}^{\text{ridge}} = (X'X + \lambda \operatorname{Id}_p)^{-1}X'Y.$$

- It can be proved that ridge penalization leads to a better generalization error.

Penalized Empirical Risk Minimization

- Penalized Empirical Risk Minimization is not restricted to linear models, it consists in finding a solution to

$$\hat{f}_{S,\lambda} := \operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f)$$

where :

- ▶ S is a class of predictors
 - ▶ pen is a (positive) penalization function, which keeps the (statistical) complexity of f under control,
 - ▶ λ is a tuning or smoothing parameter, that balances goodness-of-fit and penalization.
- λ can be chosen by cross validation.
 - When $f = f_\beta$ with $\beta \in \mathbb{R}^p$, $\operatorname{pen}(f_\beta) = \operatorname{pen}(\beta)$:
 - ▶ the ℓ_0 penalty $\operatorname{pen}(\beta) = \|\beta\|_0$,
 - ▶ the ridge penalty $\operatorname{pen}(\beta) = \|\beta\|_2$,
 - ▶ the lasso penalty $\operatorname{pen}(\beta) = \|\beta\|_1$.
 - (Data driven penalty, Radamacher, bootstrap ...)

Penalized linear predictors

$$\operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f_{\beta}) \quad (1)$$

with

$$S: \quad f_{\beta}^{\text{reg}}(x) = \langle x, \beta \rangle \quad \text{or} \quad f_{\beta}^{\text{cla}}(x) = \operatorname{sign}(\langle x, \beta \rangle)$$

- When $\beta_j = 0$, then feature j has no impact on the prediction.
- To control overfitting but also to improve interpretation we would like β to be sparse : many 0's in β .

Penalized linear predictors

$$\operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f_\beta) \quad (1)$$

with

$$S: \quad f_\beta^{\text{reg}}(x) = \langle x, \beta \rangle \quad \text{or} \quad f_\beta^{\text{cla}}(x) = \operatorname{sign}(\langle x, \beta \rangle)$$

- First option: the ℓ_0 penalty $\|\beta\|_0 = |\{j \in \{1, \dots, p\} \mid \beta_j \neq 0\}|$
 - ▶ corresponds to AIC, BIC and Mallows C_p and more recent ℓ_0 strategies.
 - ▶ strong mathematical guarantees.
 - ▶ ... but NP hard problem because we have to consider and find the solution for all the possible supports for β .

Penalized linear predictors

$$\operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f_\beta) \quad (1)$$

with

$$S: \quad f_\beta^{\text{reg}}(x) = \langle x, \beta \rangle \quad \text{or} \quad f_\beta^{\text{cla}}(x) = \operatorname{sign}(\langle x, \beta \rangle)$$

- The ridge (ℓ_2) penalty $\|\beta\|_2 = \sqrt{\sum_{j=1}^p \beta_j^2}$

► Lagrange duality:

$$(1) \Leftrightarrow \text{minimize } \sum_{i=1 \dots n} \ell(Y, f(X_i)) \text{ under } \|\beta\|_2 \leq C_\lambda.$$

Question: Draw the ℓ_2 unit ball in \mathbb{R}^2 . How does it intersect with the level sets of the empirical risk ?

► No strong sparsity induced !

Penalized linear predictors

$$\operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f_\beta) \quad (1)$$

with

$$S: \quad f_\beta^{\text{reg}}(x) = \langle x, \beta \rangle \quad \text{or} \quad f_\beta^{\text{cla}}(x) = \operatorname{sign}(\langle x, \beta \rangle)$$

- The Lasso (ℓ_1) penalty $\|\beta\|_1 = \sum_{j=1}^p |\beta_j|$
 - ▶ ℓ_1 appears as a convex proxy of the ℓ_0 .
 - ▶ Lagrange duality :

$$(1) \Leftrightarrow \text{minimize } \sum_{i=1 \dots n} \ell(Y, f(X_i)) \text{ under } \|\beta\|_1 \leq C_\lambda.$$

Question: Draw the ℓ_1 unit ball in \mathbb{R}^2 . How does it intersect with the level sets of the empirical risk ?

- ▶ Sparsity induced. Strong mathematical guarantees.
- ▶ E.g standard Lasso (regression with quadratic loss), logistic Lasso (classification with logistic loss), SVM lasso (hinge loss) ...

Penalized linear predictors

$$\operatorname{argmin}_{f \in S} \sum_{i=1 \dots n} \ell(Y, f(X_i)) + \lambda \operatorname{pen}(f_\beta) \quad (1)$$

with

$$S : \quad f_\beta^{\text{reg}}(x) = \langle x, \beta \rangle \quad \text{or} \quad f_\beta^{\text{cla}}(x) = \operatorname{sign}(\langle x, \beta \rangle)$$

- Other options:

- ▶ elastic-net : sum of lasso and ridge,
- ▶ group lasso : select variables by blocks,
- ▶ Fused lasso : penalizes the ℓ_1 norm of both the coefficients and their successive differences,
- ▶ mixed strategies.

Penalty shapes





FIGURE 3.11. Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$, respectively, while the red ellipses are the contours of the least squares error function.


[From The Elements of Stat Learning]

Regularization path: Ridge



[From The Elements of Stat Learning]

Regularization path: Lasso



[From The Elements of Stat Learning]

Outline

- 1 Overfitting and Bias Variance Tradeoff
- 2 Estimating the generalization error by cross validation
- 3 Regularization, model selection and penalization
- 4 Feature engineering**

Normalization

- Contrary to the standard least square estimator, ridge and lasso estimators are not scale invariant, even for the quadratic loss.
- We need to normalize the variables :
 - ▶ center, include an intercept, standardize
 - ▶ min-max scaling
 - ▶ binarize the variables
- More generally, these preliminary steps are recommended for most machine learning procedures.

Feature extraction and feature generation

- Create new variables to improve the prediction performances and use model selection / regularization methods to control the statistical complexity.
- Recipe / list of basic ideas :
 - ▶ Transform variables to obtain a better (simpler) relationship with the output : log, exponential, quadratic, sin, cos ..transformations
 - ▶ Introduce interaction variables
 - ▶ If there exists a model for Y (e.g. physical or biological model), including the corresponding variables will help the machine learning approach.
 - ▶ Times series : min / max / min-max / derivatives / spectral pattern
 - ▶ Decomposition and projection : PCA, splines, Fourier, wavelets ...
 - ▶ Images : specific pattern
- The python library [featuretools](#) generates many classical features, see for instance this [blog](#).
- Take home message: Feature engineering is generally a critical step for machine learning !